

Data Question 1: An Exploration of UN data

Guided Practice:

1. Download two CSV files and place them in the `data` folder of your local Data Question 1 repository:
 - a. Gross Domestic Product (GDP) per capita:
http://data.un.org/Data.aspx?d=WDI&f=Indicator_Code%3aNY.GDP.PCAP.PP.KD
 - **DO NOT APPLY ANY FILTERS**
 - rename the file to `gdp_per_capita.csv`
 - open it with a text editor (not Excel) and take a look
 - b. Percentage of Individuals using the Internet:
<http://data.un.org/Data.aspx?d=ITU&f=ind1Code%3aI99H>
 - **DO NOT APPLY ANY FILTERS**
 - rename the file to `internet_use.csv`
 - open it with a text editor (not Excel) and take a look
2. Create a Jupyter Notebook in the `notebooks` folder and name it `UN_Data_Exploration`.
 - You are likely to get errors along the way. When you do, read the errors to try to understand what is happening and how to correct it.
 - Use markdown cells to record your answers to any questions asked in this exercise. On the menu bar, you can toggle the cell type from 'Code' to 'Markdown'. [Here](#) is a link to a cheat sheet showing the basics of styling text using Markdown.
3. In the first cell of your notebook, import the required packages with their customary aliases as follows:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
```

Keep all imports in this cell at the top of your notebook.
4. At the bottom of your imports cell, use the `%matplotlib inline` magic command so that your plots show in the notebook *without* having to call `plt.show()` every time.
5. Using the pandas `read_csv()` function, read the GDP dataset into your notebook as a DataFrame called `gdp_df`.

- Take a look at the first 10 rows.
 - Look at the last 5 rows. Do you see a problem?
 - Redo the `read_csv()` call to correct this issue - **do not** modify the original csv file.
6. Drop the 'Value Footnotes' column, and rename the remaining columns to 'Country', 'Year', and 'GDP_Per_Capita'.
 7. How many rows and columns does `gdp_df` have? What are the data types of its columns? If any of the columns are not the expected types, figure out why and fix it.
 8. Which years are represented in this dataset? Take a look at the number of observations per year. What do you notice?
 9. How many countries are represented in this dataset? Which countries are least represented in the dataset? Why do you think these countries have so few observations?
 10. Create a new dataframe by subsetting `gdp_df` to just the year 2014. Call this new dataframe `gdp_2014`.
 11. Use `.describe()` to find the summary statistics for GDP per capita in 2014.
 12. Create a histogram of GDP Per Capita numbers for 2014 (you may wish to adjust the number of bins for your histogram). How would you describe the shape of the distribution?
 13. Find the top 5 countries and bottom 5 countries by GDP per capita in 2014.
 14. Now, return to the full dataset, `gdp_df`. Pivot the data for 1990 and 2017 (using the pandas `.pivot_table()` method or another method) so that each row corresponds to a country, each column corresponds to a year, and the values in the table give the GDP_Per_Capita amount. Drop any rows that are missing values for either 1990 or 2017. Save the result to a dataframe named `gdp_pivoted`.
 15. Create a new column in `gdp_pivoted` named `Percent_Change`. This column should contain the percent change in GDP_Per_Capita from 1990 to 2017. Hint: Percent change is calculated as $100 * (\text{New Value} - \text{Old Value}) / \text{Old Value}$.
 16. How many countries experienced a negative percent change in GDP per capita from 1990 to 2017?
 17. Which country had the highest % change in GDP per capita? Create a line plot showing this country's GDP per capita for all years from 1990 to 2017. Create another showing the country with the second highest % change in GDP. How do the trends in these countries compare?
Bonus: Put both line charts on the same plot.

18. Read in `continents.csv` contained in the `data` folder into a new dataframe called `continents`. We will be using this dataframe to add a new column to our dataset.
19. Merge `gdp_df` and `continents`. Keep only the countries that appear in both data frames. Save the result back to `gdp_df`.
20. Determine the number of countries per continent. Create a bar chart showing this.
21. Create a seaborn boxplot showing GDP per capita in 2014 split out by continent. What do you notice?
22. Read in `internet_use.csv` into a DataFrame called `internet_df`. You will most likely get an error message when doing this - figure out what is going wrong and fix it. Take a look at the first and last five rows and make any corrections to your `read_csv()` call to fix this. Again, **do not** modify the original datasets.
23. Drop the Value Footnotes column and then rename the columns to 'Country', 'Year', and 'Internet_Users_Pct'.
24. How many rows and columns does this new dataset have? What are the types of its columns? Ensure that the `Internet_Users_Pct` column is a numeric data type and fix it if it is not.
25. What is the first year that has a nonzero internet users percentage reported? What is the general trend in internet users percentage over the years. Pick a visualization that you think illustrates this trend well.
26. Merge `gdp_df` and `internet_df` (on Country and Year) into a single DataFrame named `gdp_and_internet_use`. Keep only countries and years that appear in both tables.
Difficult Bonus: Do not attempt this part until you have completed all other portions of the data question. Some countries have slightly different names in the internet use and gdp dataframes. For example, the Central African Republic is "Central African Republic" in the gdp dataframe and "Central African Rep." in the internet use dataframe. Find as many instances like this as you can and resolve them so that when merging you keep the maximum number of countries possible.
27. Look at the first five rows of your new data frame to confirm it merged correctly. Also, check the last five rows to make sure the data is clean and as expected.
28. Create a new DataFrame, named `gdp_and_internet_use_2014` by extracting data for the year 2014 from `gdp_and_internet_use`. What is the mean internet users percentage in 2014? How many countries have at least 90% internet users in 2014?

29. Find the countries that had the top 3 largest GDP per capita figures for 2014. Create a seaborn FacetGrid showing the change in internet user percentage over time for these five countries. Each individual figure in the facet grid will represent a single country. What trends do you notice?
30. Create a scatter plot of Internet Use vs GDP per Capita for the year 2014. What do you notice?
31. Find the correlation between GDP per Capita and Internet Use for the year 2014. What is the meaning of this number?
32. Add a column to `gdp_and_internet_use_2014` and calculate the logarithm of GDP per capita. Find the correlation between the log of GDP per capita and internet users percentage. How does this compare to the calculation in the previous part?
33. Filter the original dataset down to just the United States for all available years. Calculate correlation between internet use and gdp per capita. Is this meaningful or useful?

Solo Exploration and Presentation:

1. Choose and download another data set from the UN data <http://data.un.org/Explorer.aspx> to merge with your data and explore. Prepare a short (< 5 minute) presentation of your findings. Report any interesting correlations you find. Include visualizations and consider adding interactivity with `ipywidgets`. This presentation can be done either in a Jupyter Notebook or using another presentation software, such as PowerPoint. (Check out [Jupyter Slides](#) if you have time. This allows you to turn your jupyter notebook into a slideshow.
2. If time allows, check out the plotly library to add additional interactivity to your plots. <https://plotly.com/python/plotly-express/>.